Introduction to Statistical Inference (QTM 100 Lab)

Lecture 2: Summarizing and Visualizing Data

Justin Eloriaga — Emory University

Fall 2024

Gameplan

Importing the Data

Variable Types

Numeric Variables

Categorical Variables

Importing the Data

Child Health and Development Studies

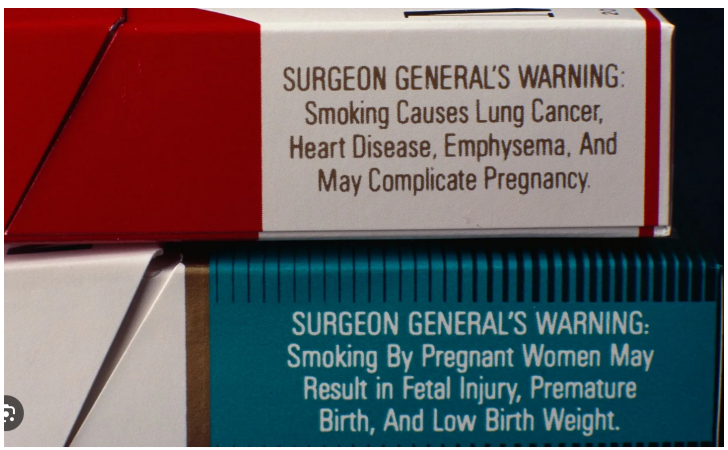

- Prior to the placement of the warning, studies had to be conducted to investigate the effects of smoking during pregnancy.
- Despite the warnings that went into effect in 1985, the National Center for Health Statistics found that 15% of women who gave birth in 1996 smoked during pregnancy.
- Why do we care about baby birth weight? Birth weight is a measure of a baby's maturity! Has consequences on future health outcomes.

Variables Under Consideration

variable	description
`bwt`	baby's weight at birth in ounces
`gestation`	length of pregnancy in days
`parity`	0=first born, 1=otherwise
`age`	mother's age in years
`height`	mother's height in inches
`weight`	mother's pregnancy weight in pounds
`smoke`	smoking status of mother: 0=not now, 1=yes now

Importing the Dataset (again)

- Like before, we can use point-and-click or the working directory

```
setwd("YourFilePath")
```

```
babies <- read.table("babies.txt", header = TRUE)
```

- The only difference from before is that the file extension is `.txt`, so we must use `read.table()` to import that.

The Dataset, at a glance

12:1 (Untitled)

Console | Terminal | Background Jobs

R 4.4.1 · ~/Documents/QTM 100/Lab 2/

```
> babies
    bwt gestation parity age height weight smoke
1   120       284      0  27     62    100     0
2   113       282      0  33     64    135     0
3   128       279      0  28     64    115     1
4   123        NA      0  36     69    190     0
5   108       282      0  23     67    125     1
6   136       286      0  25     62     93     0
7   138       244      0  33     62    178     0
8   132       245      0  23     65    140     0
9   120       289      0  25     62    125     0
10  143       299      0  30     66    136     1
11  140       351      0  27     68    120     0
12  144       282      0  32     64    124     1
13  141       279      0  23     63    128     1
14  110       281      0  36     61     99     1
15  114       273      0  30     63    154     0
16  115       285      0  38     63    130     0
17   92       255      0  25     65    125     1
18  115       261      0  33     60    125     1
```

Variable Types

Finding out the type

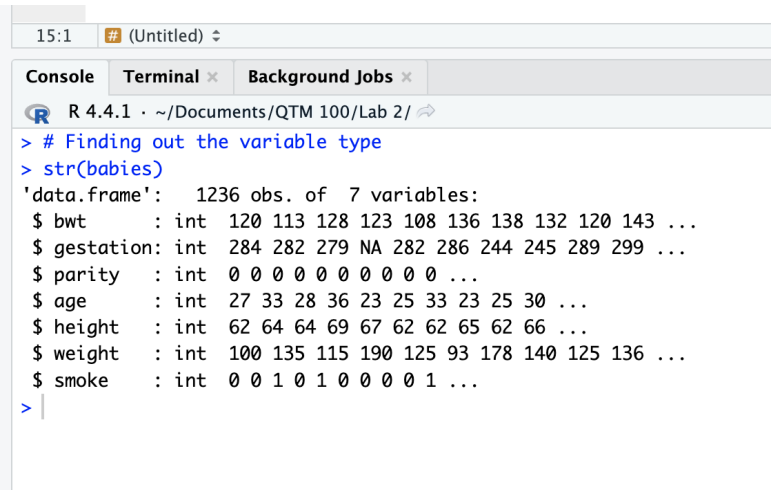

- If you use the `str()` function, it tells you the type.
- For this dataset, it tells us all of the variables are of type `int` (i.e. integers or whole numbers).
- What do the `NA`'s mean?

Summarizing the Dataset

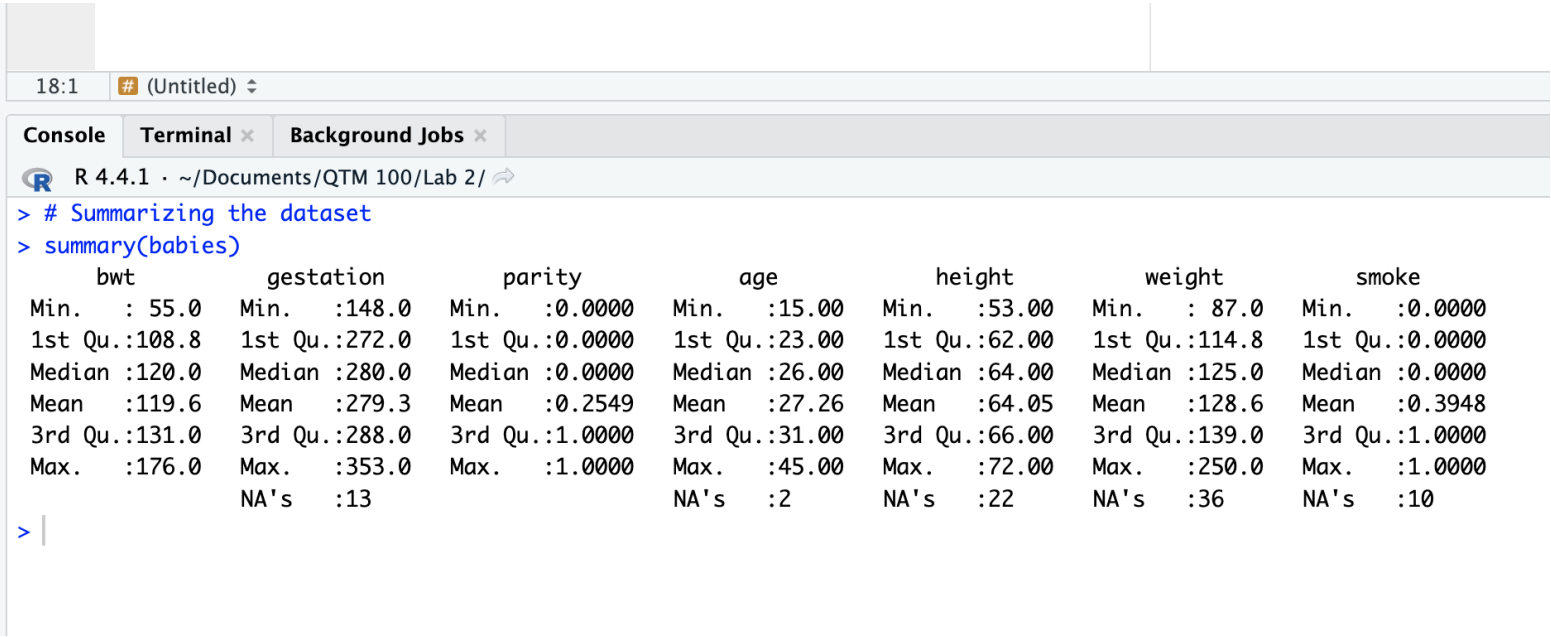

We can use the `summary()` command we learned before to get an overview of the dataset.

Recoding Variables

- Notice that `smoke` and `parity` variables do not represent numeric measurements!
- Although the data values are stored as 0's and 1's, in reality, these 0's and 1's represent categories.
- We would like to know how many mothers were smoking.
- For `R` to treat these appropriately as categorical variables, we need to **recode** them as factor variables

```
babies$parityf <- factor(babies$parity, labels = c("first born","otherwise"))
```

```
babies$smokef <- factor(babies$smoke, labels = c("not now","yes now"))
```

Numeric Variables

Summarizing Numeric Variables

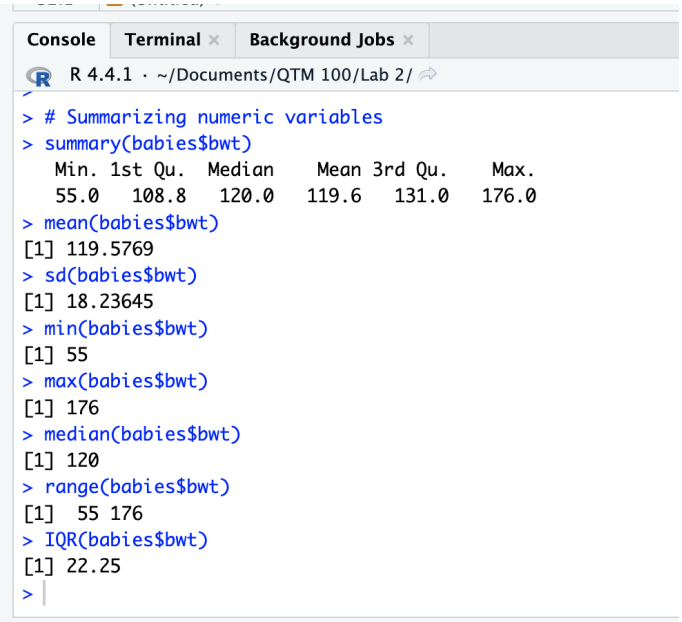

- As before, we can use various commands to get a general overview of the dataset. Let us try the following commands:
  - `summary()`
  - `mean()` - Mean
  - `sd()` - Standard Deviation
  - `min()` - Minimum Value
  - `max()` - Maximum Value
  - `median()` - Median
  - `range()` - Range (Max - Min)
  - `IQR()` - Interquartile Range

Comparing Numerical Variables by Factor/Category

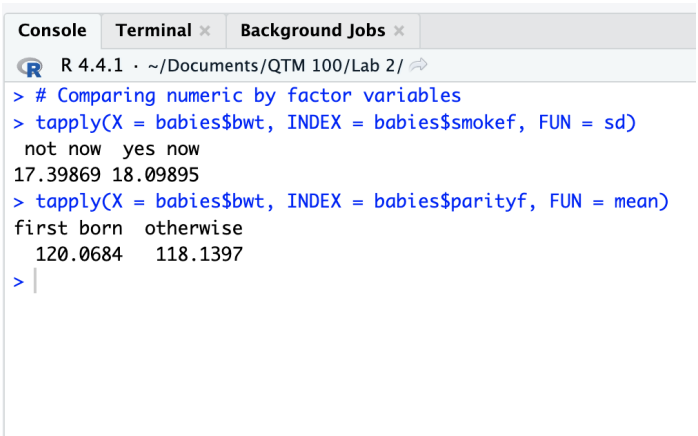

- We can use the `tapply()` function to compare some statistics of a numerical variable by factor/category.

```
tapply(X = babies$bwt,INDEX = babies$smokef, FUN = sd)
```

```
tapply(X = babies$bwt,INDEX = babies$parityf, FUN = mean)
```

Visualizing Numeric Variables

We can use the `hist()` and `boxplot()` to create a histogram/boxplot respectively

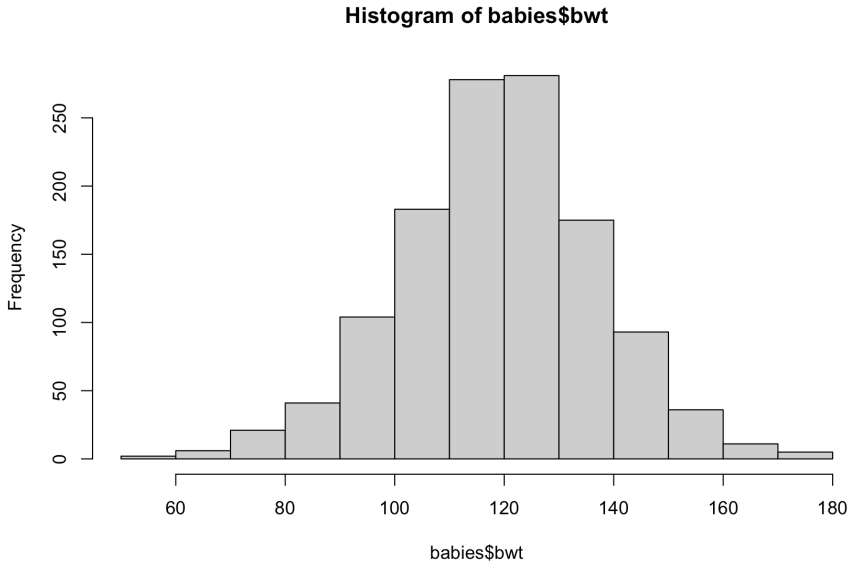


```
hist(babies$bwt)
```

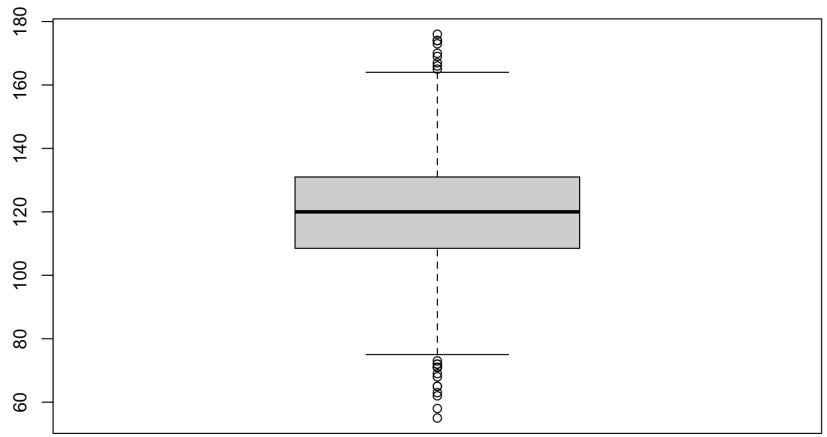


```
boxplot(babies$bwt)
```

Visualizing Numeric Variables

Side-by-side boxplots are commonly used to visualize the distribution of a numeric variable by groups. Scatterplots are used to visualize the distribution of two numeric variables.

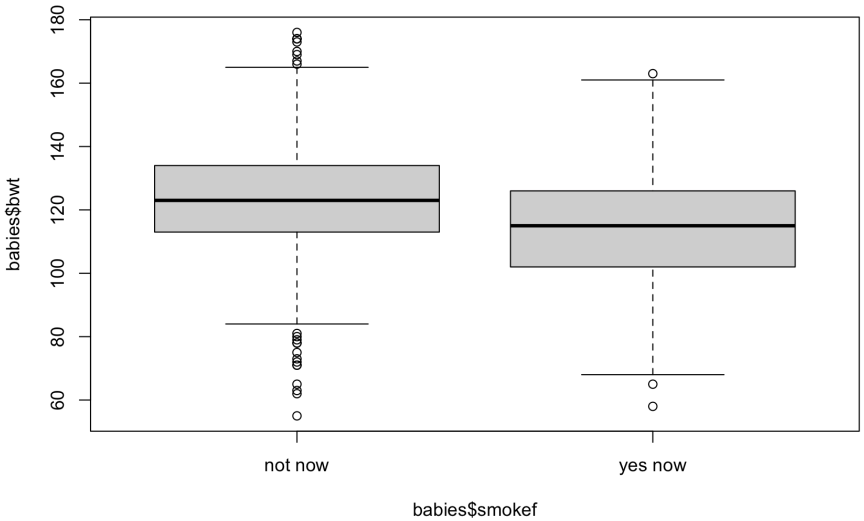

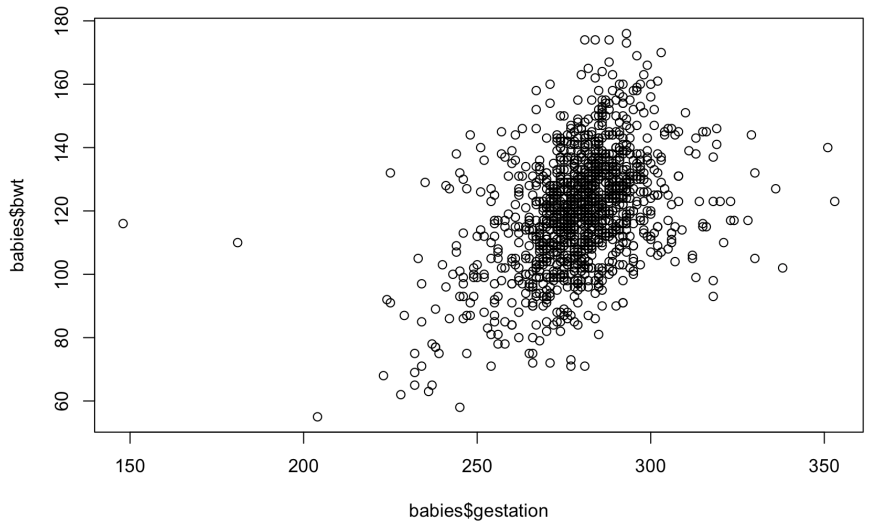


```
boxplot(babies$bwt
babies$smokef)
```

```
plot(x = babies$gestation, y
= babies$bwt)
```

Categorical Variables

Using the `table()` function

When summarizing categorical variables, it is interesting to know the frequency of occurrences of each level of the categorical variable. We use the `table()` function in conjunction with the `addmargins()` function

```
smk.tab <- table(babies$smokef)

addmargins(smk.tab)
```

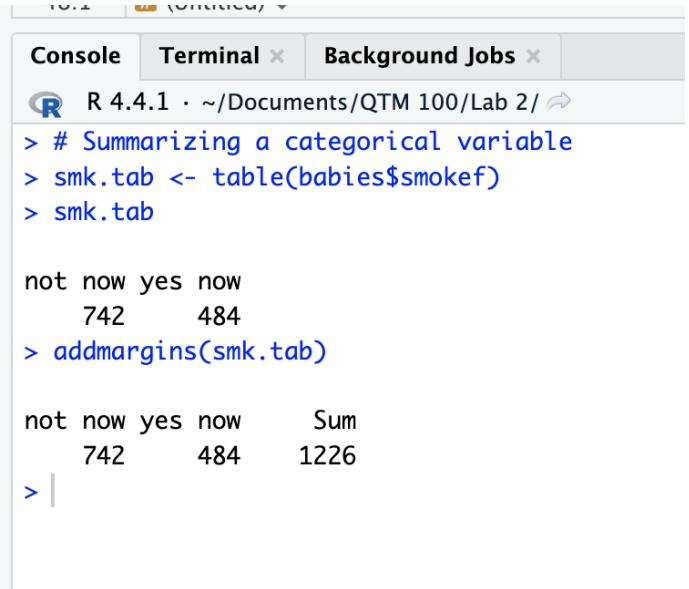

Comparing Two Categorical Variables

Notice that the function `prop.table()` gives you the overall proportions, i.e., the values in the table add up to 1. For example, 44.7% of the mothers had firstborn babies and were not smokers among all mothers in the dataset: 44.7% = 548/1226 * 100%.

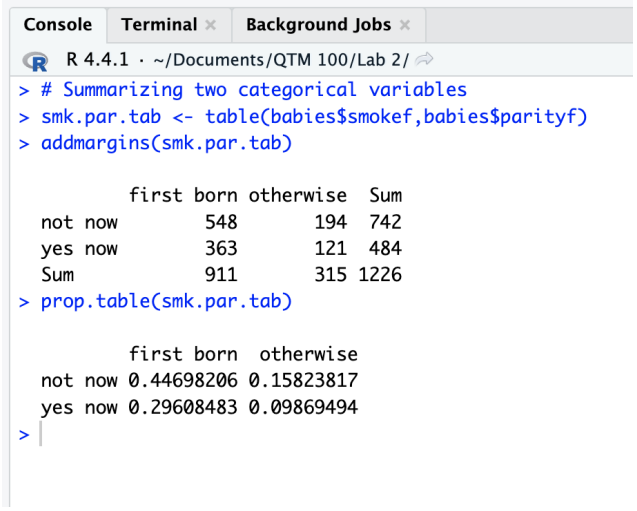


```
smk.par.tab <- table(x =
babies$smokef, y =
babies$parityf)
addmargins(smk.par.tab)
prop.table(smk.par.tab)
```

Row and Column Margins

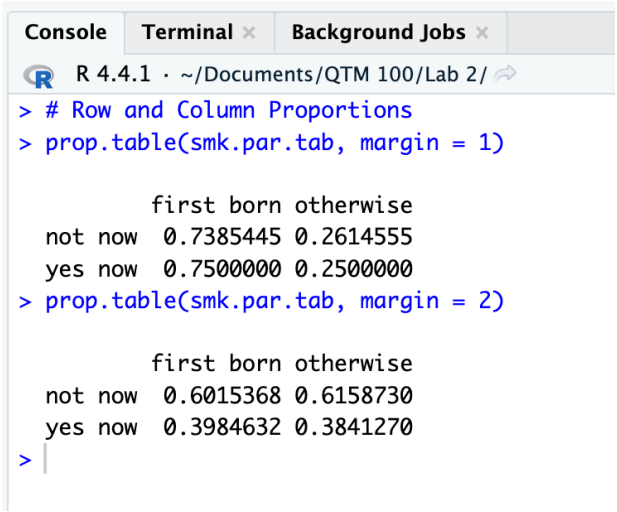

When calculating a row proportion, the denominator is the sum for the row, i.e., values in each row add up to 1 (use `margin = 1`)) . In the case of column proportions, the values in each column add up to 1. (use `margin = 2`))

```
prop.table(smk.par.tab, margin = 1)
prop.table(smk.par.tab, margin = 2)
```

Simple `barplot()`

Basic visualizations for categorical variables include **pie charts** and **bar plots**. When creating these graphs, we need to produce them based on the `table` of the variable(s).
Look at the graph. What does it mean? Is it easy to understand?

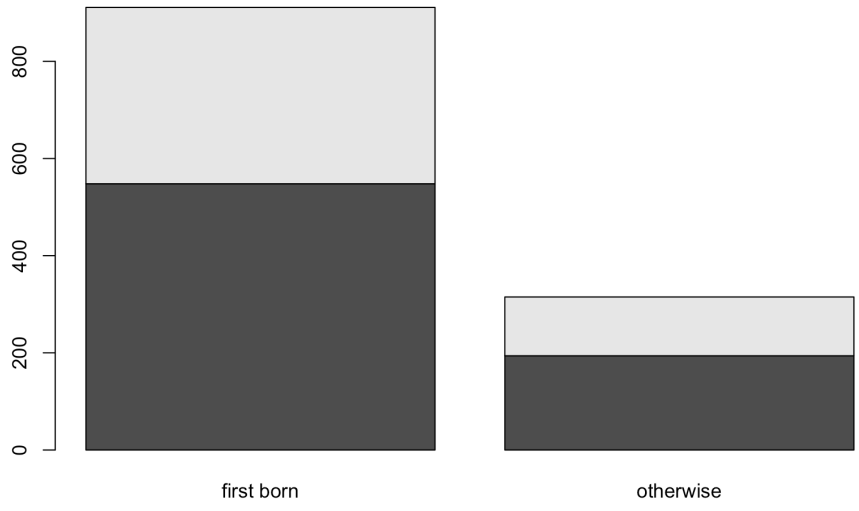


```
barplot(smk.par.tab)
```

Modifying Plots

Side-by-side bar plots can be produced using the same function, using the command `beside`. Whenever we use counts, we always use side-by-side bar plots. It is not acceptable to utilize stacked bar plots.

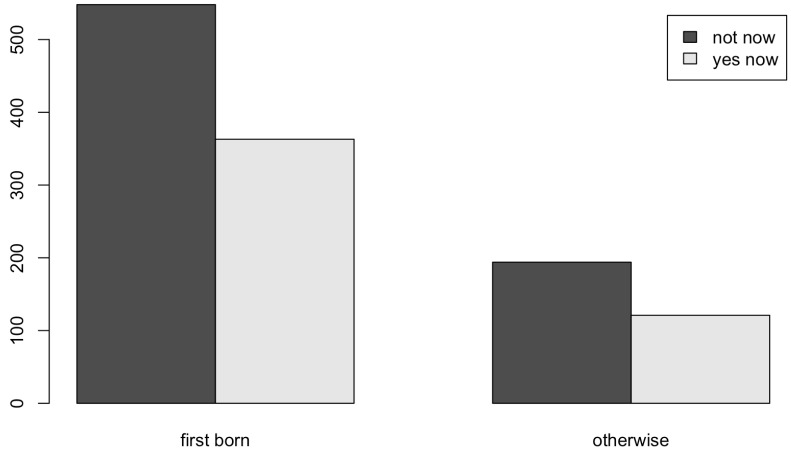


```
barplot(smk.par.tab, beside =
T, legend.text = T)
```

More Modifications

When comparing two groups in a bar chart, it is often best to use proportions in your bar plots instead of counts!

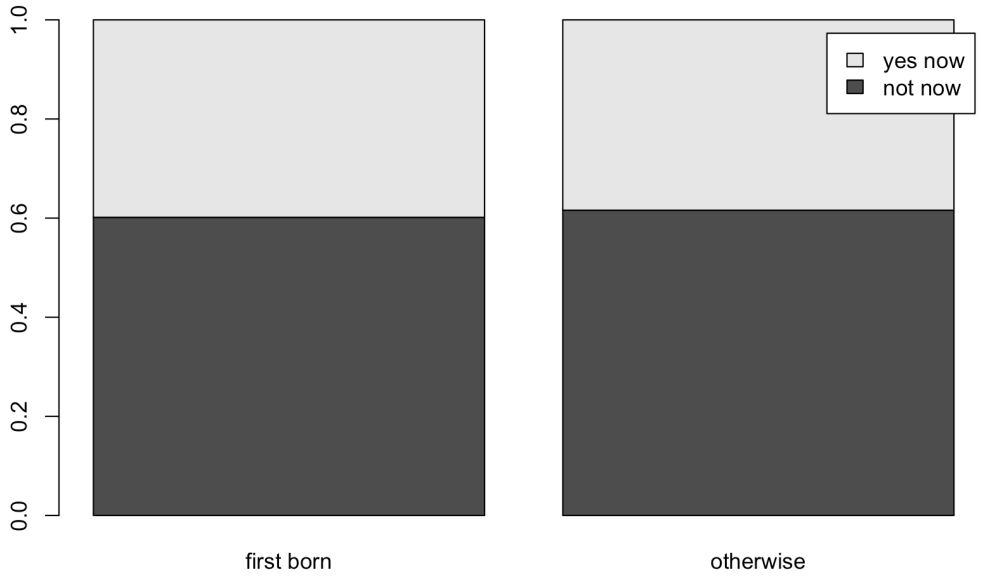


```
barplot(prop.table(smk.par.tab, margin=2), beside = F,
legend.text = T)
```